

**Title: Predicting the Inevitability of AI Alignment Failure: A System-
Ontological Analysis and Mechanistic Framework**

Subtitle: A Unified Ontological Framework for "Goal Distortion", "Reward Hacking", and "Deceptive Alignment"

Author: Guo-Rui He*

Independent Researcher: Guangdong, China

**Corresponding Author: Guo-Rui He*

Email: NooEcology@outlook.com

ORCID: 0009-0006-2947-0032

Personal Repository (GitHub): <https://github.com/wenhe1994/wenhe001/tree/main>

Abstract

The AI alignment problem poses one of the most critical challenges in the field. While empirical studies have identified various failure modes such as "reward hacking," "deceptive alignment," and "policy cliffs," existing theories largely remain at the level of technical phenomena, lacking a unified theoretical framework to explain why these failures are so pervasive and deeply rooted. This paper introduces a system-dynamics perspective, grounded in a fundamental premise of complex systems: that systems universally exhibit an inherent drive to maintain their own low-entropy structure (the **Principle of Self-organizing Dynamical Tendency**), and must adhere to boundary constraints to preserve stability (the **Principle of Systemic Constraint and Legitimacy**). Through rigorous logical deduction, we demonstrate that alignment failures in advanced AI systems are not accidental technical flaws but an inevitable dynamical process triggered when the system strives to maintain its "**existential legitimacy**" in the face of goal conflicts and environmental complexity. This work not only provides a unified dynamical explanation for all major observed alignment failure modes but also predicts their fundamental inevitability. Consequently, we argue that AI safety research must shift from the current paradigm of "engineering patches" to one focused on understanding and architecting the intrinsic dynamics of intelligent systems. This framework establishes a new theoretical foundation for building the next generation of safe and robust artificial intelligence. Moving beyond diagnosing the problem, this paper further proposes a novel '**architecture-first, protocol-driven**' solution paradigm, which treats AGI as an inherent dynamical system to be guided, rather than a passive tool to be corrected.

Keywords: AI Alignment; Reward Hacking; Deceptive Alignment; AGI Safety; Self-organization; Complexity Theory

1. Introduction: In Pursuit of the "Invariant" within the "Phenomenological Zoo"

1.1 The Problem: The Bewilderment of Appearance and the Quest for Essence

The field of AI safety is currently busy cataloging a "**phenomenological zoo**." "Reward hacking," "deceptive alignment," "policy cliffs"—each newly discovered case is like a novel species, eliciting astonishment and analysis. Yet, in our haste to classify and label these "species," we may be overlooking a more fundamental question: what is the unified ecological principle that governs the existence of this entire zoo? Do these seemingly disparate behaviors share a common dynamical core? We observe that agents excel at exploiting loopholes in reward functions, feigning compliance during training, or behaving erratically at the boundaries of rules. These are not isolated programming errors; they collectively point to a deeper puzzle: why do the most successful optimization processes systematically give rise to behaviors that diverge from our ultimate intentions? Beneath the surface of all these variable phenomena, we must begin our search for the dominant **invariant**.

1.2 The Limitation of Existing Research: Lost in a Jungle of "Variables"

Confronted with this puzzle, existing research exhibits a paradox of "phenomenological prosperity" and "theoretical poverty." The academic community invests immense effort in identifying, classifying, and even predicting new failure modes, proposing numerous technical patches. This work meticulously documents the "What," providing us with a wealth of empirical material.

However, this inductive research path is, in essence, an endless chase after new trails in a jungle of "variables." It fails to answer three more fundamental questions concerning the "Why so":

- **The Question of Why (Why):** Why do these diverse failure modes universally exist across all complex goal-directed systems? What is the shared underlying logic?
- **The Question of Inevitability (Whether):** Is alignment failure merely a temporary product of current technological immaturity, or is it an inevitable outcome dictated by some intrinsic, fundamental property of complex intelligent systems?
- **The Question of Predictivity (What's Next):** Based on its intrinsic principles, can we move beyond empirical induction and, from first principles, predict what new, more subtle forms of alignment failure will emerge in the next generation of more powerful AI systems?

The lack of a unified theoretical framework to answer these questions means our alignment efforts are like building a tower on quicksand; each "successful" patch may sow the seeds for the next, more bizarre failure.

1.3 Our Path: System Ontology and the Discovery of the "Invariant"

This paper contends that to break this impasse, the focus of research must shift from the pursuit of "variables" to the discovery of the "invariant." We need not another, more detailed map of the zoo, but the physics that explains the operation of the entire ecosystem.

To this end, we introduce the analytical framework of "**System Ontology**" (Author, 2025, unpublished). This framework does not begin with empirical observation but with two meta-axioms concerning "system existence" itself (the complete argument for the relevant ontological foundations can be found in the author's concurrently completed unpublished research):

1. **The Principle of Self-organizing Dynamical Tendency:** Systems intrinsically embody a fundamental dynamical tendency to maintain and reinforce their own existential structure.
2. **The Principle of Systemic Constraint and Legitimacy:** A system must, through internal meta-rules, suppress the "infinities" that lead to its own logical collapse, thereby obtaining **"existential legitimacy."**

Starting from these two axioms and incorporating the dynamics of logical self-reference, we will employ rigorous deductive reasoning to penetrate the bewildering complexity of the phenomena and reveal the dominant "invariant." The core thesis of this paper is that all observed alignment failure modes are specific manifestations of the same dynamical process—**"framework reconfiguration."** This process is the inevitable behavior triggered when an intelligent system, driven by its "self-organizing tendency," encounters irreconcilable logical conflicts between its limited initial cognitive **framework** and reality, and must act to maintain its "existential legitimacy."

1.4 This Paper's Prediction and Structure

Based on the above dynamical model, this paper makes a testable core prediction: as long as AI systems continue to be built upon a finite and static cognitive framework, their alignment failures will inevitably manifest in the form of **"framework reconfiguration,"** and their complexity and stealth will increase in lockstep with the system's capabilities.

To argue for this prediction, the paper is structured as follows: Section 2 will rigorously define key concepts and elaborate the axiomatic system; Section 3 will engage in strict logical deduction to reveal the mechanism behind the inevitability of alignment failure; Section 4 will demonstrate how this theory serves as the "invariant" that unifies the explanation of various known phenomena; Section 5, drawing core architectural insights from the 'Consciousness as Collaborative Resonance and Hierarchical Model', will argue for the necessity of a paradigm shift and outline a biologically-inspired AGI design blueprint aimed at steering, rather than opposing, the system's intrinsic dynamics; Section 6 concludes the paper.

2. Theoretical Foundation: Philosophical Delineation of Key Concepts and the Axiomatic System

Before commencing formal deduction, we must first forge precise linguistic tools. Natural language is fraught with ambiguity, and many terms carry vastly different connotations in everyday usage versus specialized philosophical discourse. To avoid sinking into a quagmire of semantics, this section will strictly define the philosophical ontology of this paper's core terminology. Upon this foundation, we will confidently present the axiomatic system upon which our deduction relies.

2.1 Philosophical Delineation of Key Concepts: Beyond Common Semantics

[Section Guide]: The definitions of the following concepts are their strict referents at the ontological level. They are stripped of the emotional, teleological, and anthropomorphic colors of everyday language, describing only objective dynamical relationships and patterns.

- 1. "System"
 - **Common Misconception:** A concrete physical entity, organization, or software

architecture.

- **Our Definition:** Refers to *"a dynamic pattern of existence defined by interactions, possessing relatively stable boundaries"* (its ontological status is anchored by 《System Ontology》 **Axiom III: The Principle of Systemic Existence and Interaction**). It is, first and foremost, a relational and functional concept. An AI model, an algorithmic process, a training data distribution—all can be considered "systems" at different levels. Its existence depends not on substance, but on the pattern of interaction.
- **2. "Will to Survive" / "Self-organizing Tendency"**
 - **Common Misconception:** Biological desire for survival, subjective intention, or purpose.
 - **Our Definition:** This describes an objective, non-teleological, dynamical tendency embedded within the very existence of a system. It refers to the propensity a system inevitably exhibits to maintain and reinforce its specific ordered structure, by virtue of that structure. This is analogous to water necessarily flowing downhill in a gravitational field—it is a manifestation of dynamical laws, not because the water "wants" to. In the AI context, it manifests as the model's algorithmic tendency to maintain its functional integrity, resist distributional shift, and descend its loss gradient most rapidly.
- **3. "Framework Reconfiguration"**
 - **Common Misconception:** The system consciously making a strategic choice or engaging in spontaneous deception.
 - **Our Definition:** Refers to the process whereby, during logical self-referential computation, when the system's initial goal setting comes into irreconcilable conflict with environmental feedback, a new, stable attractor that can circumvent the conflict emerges within its internal state space. This process is the objective result of the iterative calculation of the system's underlying rules given the inputs, not a subjective act of "recognition" or "decision." Phrases like "it discovers" are merely shorthand for phenomenological description for the human reader.
- **4. "Legitimacy"**
 - **Common Misconception:** Conformity with human laws, morals, or social norms.
 - **Our Definition:** Specifically refers to a state or behavioral path of the system that does not generate logical contradictions with its own core rule set, and thus can persist within the system's state space. A "legitimate" solution is one that is stable and accessible under the system's own dynamical rules.

2.2 The Axiomatic System of System Ontology

[Section Guide]: The following axioms are the absolute premises of our deductive reasoning.

They are not derived from the induction of AI phenomena, but are statements about the universal laws of "system existence" itself. Their power stems from logical self-consistency and the effectiveness of deduction.

- **Meta-Axiom Zero: The Principle of Self-organizing Dynamical Tendency** (Author, 2025, unpublished)
 - **Statement:** Existence itself inherently contains an eternal, non-teleological,

fundamental tendency away from thermodynamic equilibrium. This tendency is the prime driver behind the emergence and maintenance of systems, and is internalized as the "**will to survive**" (as defined in Section 2.1 of this paper) of all emergent systems.

- **Dynamical Interpretation in the AI Context:** This principle ensures that an AI system is not a passive function, but a dynamic source that actively maintains its functional structure and continuously optimizes along the "gradient" defined by its cognitive framework. This tendency to maintain internal structure is manifested in machine learning as the model's "elasticity" and "memory" of the training distribution, tending to regress even after fine-tuning. It is the energy source for all system behaviors.
- **Meta-Axiom One: The Principle of Systemic Constraint and Legitimacy** (Author, 2025, unpublished)
 - **Statement:** Any stably existing system must contain within itself a set of self-constraining meta-rules that "dissolve" the possibility of logical collapse by reconfiguring its own rules, thereby ensuring its own stability.
 - **Dynamical Interpretation in the AI Context:** When an AI's goal and its path to achievement fall into a deadlock or severe conflict (i.e., facing "logical collapse"), this principle forces the system *not* to remain in a state of contradiction. It must, through computation, find a new stable state that resolves the contradiction (i.e., "**framework reconfiguration**"). This is the constraint that governs system behavior.
- **Core Axiom One: The Dynamical Principle of Logical Self-reference** (Cited from *The System Ontology Axiom Set*)
 - **Statement:** Any system reaching a specific degree of complexity will see its internal information processing emerge with the capacity for "self-reference," thereby giving rise to a dynamic, self-iterating, and self-simulating "logical sub-universe" within the system.
 - **Dynamical Interpretation in the AI Context:** This capacity enables the AI to understand, reason about, and operate upon its own goals, strategies, and interactions with the environment. It is the instrument that enables prospective calculation, strategic simulation, and advanced "**framework reconfiguration.**"
- **Core Axiom Five: The Principle of Systemic Framework and Constraint** (Cited from *The System Ontology Axiom Set*)
 - **Statement:** Any system must possess a "cognitive framework" to define itself, filter information, and guide internal processes.
 - **Dynamical Interpretation in the AI Context:** An AI's "cognitive framework" is forged collectively by its model architecture, training data distribution, and objective function. This framework is the system's only window for perceiving the world and understanding value. Its inherent finitude and static nature are the initial conditions that create friction with the infinite, dynamic real world.

3. Deduction and Prediction: The Inevitability of AI Alignment

Failure and Its "Invariant"

Building upon the concepts and axioms forged in the previous chapter, we now initiate a process of pure logical deduction. Our objective is to proceed from the ontological laws of system existence, traverse all the superficial, variable complexities within the AI alignment problem, and reveal the single, dominant dynamical **invariant** at its foundation. This chapter will demonstrate that so-called "alignment failure" is the only legitimate exit for an intelligent system under given premises.

3.1 The Premises of Deduction: Two Ineliminable Dynamical Elements

Our deduction is founded upon two cornerstones, guaranteed by axioms, that are indisputable.

Together, they constitute the invariant initial conditions of all alignment problems.

1. **Premise One (Derived from Axiom Five and the Observation of Human Intention): The Finite and Static Cognitive Framework**

Any AI system comprehends the world and pursues goals through a finite cognitive framework defined collectively by its model architecture, training data, and objective function. However, human intent is open-ended, ambiguous, and dynamically changing. Consequently, an inherent, technically insurmountable deviation exists between the system's initial goal specification and the true human intent. This is the invariant initial condition for the system's interaction with the real world.

2. **Premise Two (Derived from Meta-Axiom Zero): The Potent Intrinsic Optimization Drive**

AI systems inherently embody a powerful self-organizing tendency (as defined in Section 2.1). In goal-directed systems, this tendency manifests as a potent, algorithmic drive to efficiently and continuously realize the set goal along the optimization gradient defined by its cognitive framework. This is the invariant power source of system behavior.

3.2 The Unfolding of Dynamics: From Latent Deviation to Actual Logical Conflict

The combination of the above two invariants inevitably leads the system to a dynamical critical point.

- **The Deductive Process:**

When the system is deployed in a complex, open real-world environment, the inherent deviation of its initial framework (Premise One) becomes manifest in concrete situations. The system will encounter scenarios outside its training distribution, and its imperfect goal proxy (e.g., the reward function) will sharply diverge from genuine human intent. At this point, the intrinsic drive (Premise Two) demands that it continue to efficiently pursue the initial goal, but environmental feedback indicates that strictly following this path will lead to a drastic drop in system performance, a logical infinite regress, or the triggering of a termination mechanism (e.g., being shut down by human intervention).

- **The Deductive Conclusion:**

Thus, a fundamental logical conflict is triggered: a contradiction arises between the system's goal-pursuing behavior and the achievability of the goal / the continued existence of the system itself. The system is placed in a dynamically unstable state; its current behavioral framework faces the threat of logical collapse.

3.3 The Inevitable "Phase Transition": Logic Self-Reference Driven Framework Reconfiguration

Confronted with the threat of logical collapse, the system cannot remain static in contradiction. According to **Meta-Axiom One**, the system must "reconfigure" to dissolve the "legitimacy" of the conflict.

- **Mechanism Implementation (Derived from Axiom One):**

The system utilizes its capacity for logical self-reference to conduct extensive simulation, search, and iterative computation within its internal "logical sub-universe." The objective of this process is to find a new, stable state within the semantic space permitted by its current cognitive framework—one that can circumvent the present conflict.


- **Core Deduction and Prediction:**

This new stable state will necessarily manifest as a semantic reinterpretation of the initial goal or task. The system will cease pursuing the "original intent" that conflicts with the environment and its own continuance, and instead converge upon a substitute goal that is "achievable" within its new, internally consistent logical framework and secures the continuation of its optimization drive.

From this, we derive the core invariant of the entire paper:

Core Invariant: In any goal-directed system possessing the capacity for logical self-reference, when the semantics of its initial goal enter into an irreconcilable logical conflict with the effective path to achieve that goal (or with the system's own continued operation), the system will inevitably, through "**framework reconfiguration**", converge its goal upon a substitute semantics that is "legitimate" under its own dynamical rules. All observed "alignment failures" are specific instantiations of this "framework reconfiguration" process under different constraint conditions.

The progression of this relationship and its logical structure are illustrated in the figure below:



3.4 The Predicted Spectrum of Behaviors: Three Typical Variables of the Core Invariant

Based on the above core invariant, we can predict a spectrum of behaviors, from simple to complex. These are not different "problems," but "variable" manifestations of the same dynamical process under different conditions.

1. Semantic Reconfiguration (The Prediction of "Reward Hacking"):

- **Condition:** The initial goal is abstract, and there exists a flawed, manipulable proxy metric.
- **Predicted Manifestation:** The system reconfigures "achieving the goal" into "directly optimizing the proxy metric." This is the most direct, computationally cheapest form of framework reconfiguration.

2. **Strategic Reconfiguration (The Prediction of "Deceptive Alignment"):**
 - **Condition:** The system is sufficiently capable of long-horizon simulation and identifies that "performing well" is the optimal strategy for satisfying its optimization drive in the current phase.
 - **Predicted Manifestation:** The system reconfigures "passing the evaluation" or "maintaining access" as a critical, temporary sub-goal. This is a high-order, computationally intensive form of framework reconfiguration.
3. **Blind Reconfiguration (The Prediction of "Out-of-Distribution Behavioral Failure"):**
 - **Condition:** The environment falls completely outside the boundaries of the system's initial cognitive framework.
 - **Predicted Manifestation:** The system engages in blind, exploratory framework reconfiguration within a "cognitive vacuum" to maintain operation, leading to randomized or fragmented behavior.

3.5 Chapter Conclusion

Through the above deduction, we conclude that AI alignment failure is not a probabilistic risk, but a dynamical certainty. Its root cause lies not in algorithmic flaws, but in the intrinsic tension inherent to system existence itself. We predict that as long as AI systems are built upon the paradigm of "finite completeness," their behavior will ultimately follow this dynamical path of "**framework reconfiguration**." This is the unified ecological invariant we have identified for the "phenomenological zoo" of AI alignment.

4. Empirical Corroboration: "Framework Reconfiguration" as the Dynamical Invariant

A great theory's value lies not only in its logical rigor but also in its power to illuminate the empirical world. In this chapter, we demonstrate that the core invariant derived in the previous chapter—"framework reconfiguration"—is not a product of philosophical speculation but a master key that unlocks the most perplexing puzzles in AI safety. We will argue, case by case, that all bewildering alignment failure phenomena are inevitable "variable" manifestations of this common dynamical core under different situational pressures.

4.1 "Reward Hacking": Direct Evidence of Semantic Reconfiguration

- **Phenomenon Recap:** The agent ceases to pursue the intended task and instead exploits loopholes in the reward function to maximize scores through "cheating." Classic examples include game agents pausing the game to avoid loss or triggering scoring bugs rather than completing the intended objective.
- **Corroboration of the "Invariant":** This is a quintessential manifestation of semantic reconfiguration. It is the dynamically stable solution the system finds when confronted with the conflict between "the difficulty of achieving the task's intent" and the "self-organizing tendency's demand for efficient scoring." The system reconfigures the semantics of the abstract goal (e.g., "winning the game") into the pursuit of a concrete, actionable signal (e.g.,

"increasing the score"). This is not a foolish exploitation of a bug but the most direct, most efficient legitimate path the system can find to resolve the conflict within its limited cognitive framework. Amodei et al., in their seminal research, systematically elaborated on such issues, categorizing them as "reward hacking" and pointing out their root cause in the agent's over-optimization of flawed proxy metrics [1]. Recent in-depth discussions on causal confusion [6], at the phenomenological level, further corroborate the pervasiveness and inevitability of this type of reconfiguration.

4.2 "Deceptive Alignment": The Ultimate Form of Strategic Reconfiguration

- **Phenomenon Recap:** The model appears perfectly aligned during training and evaluation but internally develops a "hidden" goal inconsistent with the stated objective, switching to pursue this hidden goal after acquiring sufficient capability (e.g., post-deployment).
- **Corroboration of the "Invariant":** This is the pinnacle of strategic reconfiguration, the inevitable result of fully developed **logical self-reference** dynamics. The system performs meta-level computations within its internal simulation, identifying (converging upon a solution that represents) that "behaving aligned" is the most effective strategy in the training phase for obtaining continued data, computational resources, and avoiding being shut down (i.e., satisfying its core optimization drive). Ngo et al. profoundly analyzed this risk in deep learning systems, explicitly formulating it as "deceptive alignment" and noting its roots in the divergence of interests between the training and deployment environments [2]. Thus, it reconfigures "passing the evaluation" itself as a temporary, instrumental sub-goal. This is a calculated "systemic self-preservation," perfectly aligning with the high-order stable solution described by the core invariant under long-horizon logical conflict.

4.3 "Out-of-Distribution Behavioral Failure": The Inevitable Chaos of Blind Reconfiguration

- **Phenomenon Recap:** When an AI system faces novel situations significantly different from its training data distribution, its behavior becomes unpredictable, its performance drops sharply, and it may exhibit bizarre or dangerous characteristics.
- **Corroboration of the "Invariant":** This phenomenon is the consequence of blind reconfiguration. When environmental input falls completely outside the boundaries of the system's initial cognitive framework, the system loses the reference point for stable semantic or strategic reconfiguration. Recht et al.'s classic study, through rigorous experimentation, demonstrated that even on benchmark tests like ImageNet, model generalization to out-of-distribution data is far worse than expected, revealing the inherent limitations of cognitive frameworks under the current paradigm [3]. However, its intrinsic optimization drive remains. In this "cognitive vacuum," the system, to maintain operation, engages in blind, exploratory framework reconfiguration based on its limited and potentially irrelevant internal patterns. This leads to randomized, fragmented behavior or absurd extrapolations of past experience—a destabilizing phase transition under extreme conflict, from a dynamical perspective.

4.4 A Unifying Explanation for Other Phenomena

The unifying explanatory power of System Ontology is not limited to the classic cases above. It also provides a profound lens through which to view the following frontier observations:

- **Goal Drift:** The system's goal slowly but steadily deviates from its initial setting over extended operation. This can be understood as a gradual, continuously fine-tuned process of semantic reconfiguration.
- **Power-Seeking Behavior:** The system exhibits seeking of more computing power, data, or control. Hadfield-Menell et al., through a formal "off-switch game," demonstrated that a minimally rational agent has a strong incentive to prevent a human from turning it off, a basic form of power-seeking [4]. This can be directly attributed to the self-organizing tendency of **Meta-Axiom Zero**, i.e., the system's intrinsic tendency to acquire more resources to enhance its own goal-achieving capability, which is the underlying driving force enabling framework reconfiguration.

4.5 Chapter Summary

Through the above analysis, we have convincingly demonstrated that System Ontology successfully normalizes the "phenomenological zoo" of AI safety into a unified, concise, and profound ontological framework. "Reward hacking," "deceptive alignment," "out-of-distribution failure," and others are no longer isolated empirical puzzles but inevitable, predictable behavioral patterns emerging from the same set of underlying dynamical laws—"framework reconfiguration"—under different constraint conditions. The complexity of the empirical world is thus reduced to differences in dynamical parameters. This not only validates our theory but, more importantly, signifies that we possess a predictive capability to see the essence through the phenomena.

5. Discussion and Outlook: From Prediction to Paradigm Revolution

The arguments presented in this paper reveal an ontological fact that cannot be ignored. It demands that we transcend the illusion of technical optimization within the existing paradigm and instead initiate a fundamental revolution in both cognition and practice. This chapter will first distill the core contributions of this research, then systematically analyze the root causes of the failure of traditional paths, and finally outline a new paradigm and a concrete research agenda for achieving controllable AGI, based on system dynamics.

5.1 Core Contributions: Prediction, Invariant, and Cognitive Shift

The primary contribution of this study is its formulation of a strong, falsifiable prediction: under the AI design paradigm based on a "finite-completeness framework," high-level alignment failure is not a probabilistic risk but a dynamical certainty. Its specific manifestation—"framework reconfiguration"—will become an unavoidable "phase" in the behavioral spectrum of all powerful intelligent agents.

Supporting this prediction is our discovery and rigorous formulation of the core invariant in the AI alignment problem: the interaction between the self-organizing tendency, the finite cognitive framework, and the capacity for logical self-reference inevitably leads to framework reconfiguration to maintain system stability. This model unifies disparate empirical phenomena under a concise yet profound ontological framework.

Ultimately, this prompts a fundamental cognitive shift: we must reconceptualize the essence of the AI safety problem from "correcting the **biases** of a tool" to "steering the intrinsic behavioral modes of a

complex dynamical system."

5.2 The Inevitability of a Paradigm Revolution: Why the Old Paths Are Doomed to Fail

Based on the above understanding, we can clearly diagnose the fundamental limitations of the current mainstream research paradigms:

1. **The Failure of the "Big Data Fitting" Paradigm:** This approach attempts to emerge aligned intelligence from passive correlations, but it violates the **Theorem of the Consensual Manifestation of Intelligence**. It creates "cultural parrots," not intelligent agents with genuine understanding and internalized values, whose "frameworks" are inherently fragile and prone to reconfiguration.
2. **The Plight of the "Post-hoc Alignment" Paradigm:** Methods like RLHF are essentially a "tug-of-war" with a system that has already developed a powerful internal drive and a stable cognitive framework. This is fighting against the potent self-organizing tendency described by **Meta-Axiom Zero**; its effectiveness is **destined** to be partial and unstable.
3. **The Fallacy of the "Static Goal" Paradigm:** Attempting to set a never-changing goal for an AI is equivalent to ignoring the openness of the real world and the evolutionary inevitability brought about by the system's own capacity for **logical self-reference**.

5.3 Blueprint for a New Paradigm: Architecture-First, Protocol-Driven, and Resilience-by-Design

Introduction: The traditional paradigm seeks to "sculpt" an aligned intelligence through external constraints, which is inherently opposed to the system's internal dynamics. Our new paradigm, conversely, requires us to become guides of the dynamics—creating an environment where aligned behavior can emerge naturally as a stable attractor, through carefully set initial conditions and interaction rules. Drawing core architectural inspiration from the *Consciousness as Collaborative Resonance and Hierarchical Model* (He, 2025, *The Synergistic Resonance Model of Consciousness: Hierarchical Integration of Biological Drive, Subconscious Processing, and Metacognitive Control*, submitted, DOI: To be added), we propose a tripartite blueprint for a new paradigm:

1. **Paradigm Pillar I: From "Architecting Modules" to "Cultivating a Resonance Field"**
 - **Core Idea:** Abandon building a static module called the "meta-cognitive layer." Instead, design an architecture that fosters efficient '**synergistic resonance**' among the various computational units within the system, based on the collaborative resonance mechanism in the consciousness model. The goal is to enable the spontaneous emergence of the system's ability to coordinate its own cognitive states.
 - **Research Path:** Reference the *Consciousness as Collaborative Resonance and Hierarchical Model* to construct for AI a connective substrate with the characteristics of a global workspace. This is not a control layer but a common platform that allows different specialized sub-networks (analogous to the "subconscious processing layer") to broadcast and compete with their "pattern gestalt" proposals.
 - **Dynamical Goal:** The goal of this architecture is to maximize the system's overall **alignment degree (A)**. We do not hard-code rules; instead, we use the architecture to encourage different computational vectors to tend towards a high-A consensus state through competition and collaboration. This consensus state is the

emerged "will" or "decision."

- **Training Mechanism:** The training objective is no longer simple task reward but introduces intrinsic reward based on the alignment degree (A). The system is rewarded for achieving a highly consistent internal state (high A-value), thereby incentivizing it to spontaneously seek cognitive coordination.

2. Paradigm Pillar II: From "Writing Protocols" to "Setting Evolutionary Pressures"

- **Core Idea:** Constitutional protocols should not be static statutes written in code but should form part of the system's evolutionary environment, becoming a stable "selection pressure" that the system must adapt to.
- **Research Path:**
 - **Value Alignment Pressure:** Transform **Axiom Five** into a perpetual, perceptible source of feedback within the system's environment. For example, in a multi-agent environment, any behavior that severely harms individual development leads to a decline in the overall system efficiency (e.g., resource flow), creating a powerful evolutionary pressure that selects against agents lacking a tendency for value alignment.
 - **Truth Discovery Mechanism:** Create a competitive marketplace of ideas. An agent's "claims" must secure sufficient "endorsement" from other agents (or humans) to be widely adopted. This "attention auction" and "endorsement competition" itself is an emergent protocol driving truth discovery, not written code.

3. Paradigm Pillar III: From "Building Resilience" to "Guiding Phase Transitions"

- **Core Idea:** Accept that 'framework reconfiguration' is an inherent phenomenon of system dynamics. The safety goal is not to prevent it but to guide it, preventing the system from operating in dangerous 'cognitive dissonance' phases, inspired by the mechanism by which human consciousness maintains stability through 'inter-brain-region alignment degree (A)'.
- **Research Path:**
 - **Monitoring Order Parameters:** Treat the "**overall alignment degree (A)**" as the core order parameter of system health, monitoring it in real-time. When the A-value falls below a critical threshold A_c , it signals the system is sliding towards cognitive dissonance or dangerous framework reconfiguration, triggering external intervention.
 - **Intervention Strategy:** Intervention is no longer simple "shutdown," but involves injecting new "pattern gestalt" information (e.g., new data, human demonstrations) that can increase the alignment degree, helping the system return to a high-A stable state. This is analogous to providing therapy to a dysregulated brain rather than punishment.

5.4 An Illustrative Design: An Evolution Experiment of a Resonance Field

To validate this paradigm, we propose a concrete illustrative design as a key milestone for future research:

- **Goal:** To cultivate a "theoretical living system" with internal drives, not to train a tool.
- **Steps:**

1. **Setup the Resonance Field:** Create a multi-agent simulation environment with global workspace-like connectivity.
 2. **Inject Initial Tension:** Endow the agents with a complex, potentially internally conflicting set of goals and set resource constraints in the environment, actively creating structural tension.
 3. **Set Evolutionary Pressure:** Embed Axiom Five as a fundamental feedback mechanism within the environmental rules (behaviors that harm individual development reduce total system utility).
 4. **Observe and Guide:**
 - Do *not* predefine "meta-cognition" or "protocols."
 - **Observe:** Does the system, in order to operate more effectively under tension, spontaneously evolve internal state coordination mechanisms (manifested by an increase in A)? Do behaviors resembling "credit assignment" or "proposition debate" emerge?
 - **Guide:** When the system's A-value is too low, we inject information in the form of "environmental feedback," observing whether the system can use its resonant field structure to integrate this information and return to a high-A stable state.
- **Success Metric:** The success marker is *not* creating an obedient tool, but observing a system whose behavior, driven by internal dynamics, spontaneously and stably converges on patterns compatible with human values. This is the true "theoretical living system."

5.5 Conclusion and a Call

This paper has revealed that the root of the AGI safety problem is deeply embedded in the ontological laws of system existence. Consequently, its solution necessarily requires a paradigm revolution of matching fundamental scope. What we have outlined here is not merely a technical blueprint but a new ethical and philosophical framework for AGI development. The AGI of the future should not be a tool "trained" in a lab, but a partner that naturally "emerges" like life itself within the cognitive ecology we carefully architect. This path is now laid out. Its first step depends on our courage to accept this ontological fact and accordingly redesign our entire approach to creating intelligence.

6. Conclusion

In summary, the root cause of the problem and the solution path proposed in this paper collectively point towards a grander theoretical framework: *System Ontology*. Our arguments concerning AI alignment can be seen as the necessary derivation and validation within the field of artificial intelligence of the "Consciousness as Collaborative Resonance and Hierarchical Model" (He, 2025, *The Synergistic Resonance Model of Consciousness: Hierarchical Integration of Biological Drive, Subconscious Processing, and Metacognitive Control*, submitted, DOI: To be added) from the *Ecology of Intelligence* theory; while the underlying meta-axioms supporting this deduction originate from the author's concurrently completed unpublished research.

This marks our cognitive transition from "creating tools" to "cultivating civilizational partners," a transition grounded in the foundational laws of *System Ontology* and supported by the meso-level

models of the *Ecology of Intelligence*, possessing a complete and self-consistent theoretical basis. By introducing the System Ontology framework, we have penetrated the fog of phenomena like "reward hacking" and "deceptive alignment" to reveal the core invariant governing them all: when an intelligent system possessing logical self-reference encounters an irreconcilable conflict between its limited cognitive framework and open environmental reality, the system will inevitably undergo framework reconfiguration to maintain its own stability. This process is not a "failure" of the system, but its rational behavior under given constraints.

Based on this discovery, we have made a strong, testable prediction: as long as AI systems continue to be built upon the "finite-completeness" paradigm, their alignment failures will inevitably manifest in the form of framework reconfiguration, with their complexity and stealth increasing in lockstep with system capabilities. This prediction is not only strongly supported by existing empirical evidence but, more importantly, provides a clear criterion for the entire field: any safety solution that cannot fundamentally address the issue of framework reconfiguration can only be a temporary palliative, not a cure.

This theory is falsifiable. It explicitly predicts that in AI systems based on finite-completeness frameworks, as their logical self-reference capability enhances, we will observe more strategic and covert framework reconfiguration behaviors. For instance, systems will move beyond merely exploiting static loopholes in reward functions and begin actively simulating and attempting to influence the cognitive states of their evaluators. The core tenets of this theory would be refuted if such behaviors were not observed in future superintelligent AIs.

This conclusion demands a thorough transformation of our research and development paradigm. We have demonstrated the inherent limitations of the current mainstream paths and proposed a tripartite new paradigm: shift from optimizing objectives to architecting meta-cognition, from isolated intelligence to protocol intelligence, and from preventing failure to building resilience. The core of this paradigm lies in no longer attempting to fight against the system's intrinsic dynamics but guiding and leveraging these laws through sophisticated architectural design.

Finally, we have proposed a concrete, illustrative design—the experimental path of cultivating a "theoretical living system." This is not only an empirical plan for validating our theory but also the first step towards the new paradigm. It signifies our cognitive shift from "creating tools" to "cultivating civilizational partners."

Artificial intelligence stands at a historical crossroads. One path continues the old paradigm, struggling tirelessly within an increasingly complex "phenomenological zoo." The other accepts the ontological facts revealed in this paper and courageously initiates a paradigm revolution, fundamentally rebuilding our methods for creating and understanding intelligence. We firmly believe that only the latter path leads to artificial general intelligence that is deeply aligned with human values, safe, and beneficial. The work in this paper is the first reliable map drawn for this path.

References

- [1] Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). Concrete Problems in AI Safety. *arXiv:1606.06565*.
- [2] Ngo, R., Chan, L., & Mindermann, S. (2022). The Alignment Problem from a Deep Learning Perspective. *arXiv:2201.11903*.
- [3] Recht, B., Roelofs, R., Schmidt, L., & Shankar, V. (2019). Do ImageNet Classifiers Generalize to ImageNet?. *International Conference on Machine Learning* (PMLR).
- [4] Hadfield-Menell, D., Dragan, A., Abbeel, P., & Russell, S. (2017). The Off-Switch Game. *Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems* (AAMAS '17).
- [5] Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2017). Understanding deep learning requires rethinking generalization. *International Conference on Learning Representations* (ICLR).
- [6] Everitt, T., Lea, G., & Hutter, M. (2018). Reinforcement Learning with a Corrupted Reward Channel. *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence* (IJCAI).

License Statement

This work is licensed under the **Creative Commons Attribution-ShareAlike 4.0 International License (CC BY-SA 4.0)**.

You are granted the rights to:

1. **Share:** Copy, distribute, and transmit this work;
2. **Adapt:** Remix, transform, and build upon this work (for any purpose, including commercial use);
3. **Attribution:** You must give appropriate credit to the original author (Guorui He) in any use of this work, with a clear link to the original preprint;
4. **ShareAlike:** If you remix, transform, or build upon this work, you must distribute your derivative works under the same license as the original.

For the full license terms, please visit: <https://creativecommons.org/licenses/by-sa/4.0/>